

Problem Set 1

Gen Nishida

Handed In: September 7, 2014

1 Review Questions

1. Assume the probability of getting *head* when tossing a coin is λ .

- What is the probability of getting the first head at the $(k+1)$ -th toss?

$$(1 - \lambda)^k \lambda$$

- What is the expected number of tosses needed to get the first head?

Let k be the number of tosses needed to get the first head. Then the expected number of k is defined by

$$E[k] = \lambda \times 1 + (1 - \lambda)\lambda \times 2 + (1 - \lambda)^2 \lambda \times 3 + \dots \quad (1)$$

$$(1 - \lambda)E[k] = (1 - \lambda)\lambda + (1 - \lambda)^2 \lambda \times 2 + (1 - \lambda)^3 \lambda \times 3 + \dots \quad (2)$$

By subtracting (2) from (1), we get

$$\begin{aligned} \lambda E[k] &= \lambda + (1 - \lambda)\lambda + (1 - \lambda)^2 \lambda + (1 - \lambda)^3 \lambda + \dots \\ &= \lambda \frac{1}{1 - (1 - \lambda)} \\ &= 1. \end{aligned}$$

Thus, $E[k] = 1/\lambda$.

2. Let $f(x, y) = 3x^2 + y^2 - xy - 11x$

- What is the partial derivative of f with respect to x ($\frac{\partial f}{\partial x}$)? Find $\frac{\partial f}{\partial y}$ as well.

$$\frac{\partial f}{\partial x} = 6x - y - 11, \quad \frac{\partial f}{\partial y} = 2y - x$$

- Find a point (x, y) that minimizes f .

$$\begin{cases} 6x - y - 11 = 0 \\ 2y - x = 0 \end{cases}$$

By solving these equations, we get $(x, y) = (2, 1)$.

3. • Assume that $\omega \in \mathbb{R}^n$ and b is a scalar. A hyperplane in \mathbb{R}^n is the set $\{x : x \in \mathbb{R}^n, \omega^T x + b = 0\}$. For $n = 2$ and $n = 3$, draw on paper an example of a hyperplane. The hyperplane has its normal vector ω , and it is away from the origin by $-b/\|\omega\|$. The example of a hyperplane for $n = 2$ and $n = 3$ is in Figure 1.




Figure 1: The example of a hyperplane

- Assume we have two parallel hyperplanes: $\{x : x \in \mathbb{R}^n, \omega^T x + b_1 = 0\}$ and $\{x : x \in \mathbb{R}^n, \omega^T x + b_2 = 0\}$. What is the distance between these two hyperplanes?

$$\left| \frac{-b_1}{\|\omega\|} - \frac{-b_2}{\|\omega\|} \right| = \frac{|b_1 - b_2|}{\|\omega\|}$$

2 Basic Concepts

1. Define in one sentence: (1) training set, (2) test set, (3) validation set.

- training set

Training set is a set of data used to optimize a hypothesis function.

- test set

Test set is a set of real-world data used to measure the accuracy of the hypothesis generated through training and validation phases.

- validation set

Validation set is a set of data used to estimate the performance of the hypothesis during the training phase.

2. Can you use the validation set as a test set?

No. Since validation set is used to estimate the accuracy of the hypothesis during the training phase, the resulting hypothesis is optimized for the validation set, and it is meaningless to use the validation set as a test set in order to measure the actual performance for the real-world data.

3. Define in one sentence: overfitting

A hypothesis is said to overfit the training data if it has smaller error on the training data but loses the generalization performance and has larger error on test data.

4. True or False (and why): A learned hypothesis f has a training error e_{tr} and a testing error e_{ts} , where $e_{tr} > e_{ts}$.

- (1) can we say that f overfits to the training data?

False. Since the hypothesis f is optimized for the training data while the test data is unknown during training phase, the training error e_{tr} is smaller than e_{ts} in general, even if f is well generalized. In this case, $e_{tr} > e_{ts}$, which indicates that f is generalized very well.

- (2) Now, assume that $e_{tr} < e_{ts}$, does f overfit to the training data?

False. Since the hypothesis f is optimized for the training data while the test data is unknown during training phase, the training error e_{tr} is smaller than e_{ts} in general, even if f is well generalized. Therefore, we cannot conclude that f overfits to the training data even if $e_{tr} < e_{ts}$, unless we find another hypothesis f' which has larger error on the training data but smaller error on the test data compared to f .

3 Decision Trees

1. The "Thrill and Romance" bookstore

- What is the entropy of the target variable? (Buy)

The number of examples labeled "Buy=Y" is 7, while the number of examples labeled "Buy=N" is 4. Thus,

$$-\frac{7}{11} \log \frac{7}{11} - \frac{4}{11} \log \frac{4}{11} = 0.94566$$

- What are the attributes considered by the algorithm?

All the attributes, "Pages", "Famous Author", "Category", and "Cover Color" should be considered by the algorithm. However, for "Pages", since it is a continuous attribute, we first have to sort examples according to the values of "Pages" and check the mid-point as a possible threshold in order to discretize. The sorted values are as follows:

45(-), 50(+), 72(+), 100(-), 120(+), 142(+), 150(+), 200(-), 300(+), 350(+), 1000(-).

Thus, the possible thresholds for "Pages" are 47.5, 86, 110, 175, 250, and 675. These values are called cut points [1], and used to split the continuous space into two ranges so that the continuous values can be handled in the same manner with the discrete values.

- What is the first attribute that the algorithm will split the data on? What is its information gain?

The information gain by the split of each attribute is computed as follows.

– Pages(threshold=47.5)

$$0.94566 - \left[0 \times \frac{1}{11} - \left(-\frac{7}{10} \log \frac{7}{10} - \frac{3}{10} \log \frac{3}{10} \right) \times \frac{10}{11} \right] \\ = 0.94566 - 0.80117 = 0.14449$$

– Pages(threshold=86)

$$0.94566 - \left[\left(-\frac{2}{3} \log \frac{2}{3} - \frac{1}{3} \log \frac{1}{3} \right) \times \frac{3}{11} + \left(-\frac{5}{8} \log \frac{5}{8} - \frac{3}{8} \log \frac{3}{8} \right) \times \frac{8}{11} \right] \\ = 0.94566 - 0.94458 = 0.00108$$

– Pages(threshold=110)

$$0.94566 - \left[1.0 \times \frac{4}{11} - \left(-\frac{5}{7} \log \frac{5}{7} - \frac{2}{7} \log \frac{2}{7} \right) \times \frac{7}{11} \right] \\ = 0.94566 - 0.91289 = 0.03277$$

– Pages(threshold=175)

Same as threshold=110, which is 0.03277.

– Pages(threshold=250)

Same as threshold=86, which is 0.00108.

– Pages(threshold=675)

Same as threshold=47.5, which is 0.14449.

– Famous Author

$$0.94566 - \left[\left(-\frac{5}{7} \log \frac{5}{7} - \frac{2}{7} \log \frac{2}{7} \right) \times \frac{7}{11} + 1.0 \times \frac{4}{11} \right] \\ = 0.94566 - 0.91289 = 0.03277$$

– Category

$$0.94566 - \left[\left(-\frac{4}{5} \log \frac{4}{5} - \frac{1}{5} \log \frac{1}{5} \right) \times \frac{5}{11} + 1.0 \times \frac{6}{11} \right] \\ = 0.94566 - 0.8736 = 0.07206$$

– Cover Color

$$0.94566 - \left[\left(-\frac{6}{9} \log \frac{6}{9} - \frac{3}{9} \log \frac{3}{9} \right) \times \frac{9}{11} + 1.0 \times \frac{2}{11} \right] \\ = 0.94566 - 0.93315 = 0.01251$$

The attribute which produces the highest gain is "Pages", so we should use "Pages" as the first attribute to split, and its information gain is 0.14449 if we use threshold=47.5 or 675. Note that if we split all the values of "Pages", then we will easily be able to separate the labels with information gain = 1.0. However, this will be overfitting, and we do not want to do that.

- Due to a computer error some of the training examples attributes were deleted! Revise the decision tree training algorithm to deal with missing values in the training data.

The goal of training is to minimize the expected loss over the distribution of values of attributes. Intuitively, it is better to use this distribution to estimate the missing values, but, we do not know the distribution in advance. One option would be to assume that the missing values can be all the possible values with equal probability.

Suppose the first example in our training data does not have a value of "Famous Author". If we use "Famous Author" as the first attribute to split the tree, then this example will contribute a half to each sub-tree. In other words, the subset S_Y for which attribute "Famous Author" has value "Y" will contain 5.5 examples, while the other subset S_N will have 4.5 examples. Also, since the first example is labeled "Buy=Y", in each subset, this example contributes 0.5 to the positive proportion for computing entropy, and the information gain can be computed in the similar manner. Another option is to use the proportion of each sub-tree as a probability of having the corresponding value. If we have enough number of training data without missing values, the proportion of the sub-trees are likely to be similar to the distribution of testing data. Hence, we can estimate the missing values. In the decision tree implementation in the following section, I used the latter approach.

2. Decision Tree Implementation

I implemented the decision tree algorithm using the second option for the missing values. With regard to the continuous values, for every node, I used each cut point discussed above as a candidate of the attribute in addition to the discrete attributes, and choose the best one which produces the highest information gain. Also, the validation data was used to find the best maxDepth, and measure the performance of the resulting decision tree on the test data. Please refer to the attached source code for details.

Table 1: Results

maxDepth	Validation data	Test data
0	0.58461	0.50746
1	0.96923	0.52238
2	0.96923	0.52238
3	0.96923	0.52238
4	0.90769	0.56716
5	0.92307	0.56716
6	0.92307	0.56716
7	0.92307	0.56716
8	0.92307	0.56716
9	0.92307	0.56716

The results of my decision tree on the validation and testing data are shown in Table 1. Since the testing data is not available beforehand, we have to choose the best `maxDepth` only based on the results on the validation data. Thus, we choose *maxDepth* = 1, and we get 0.52238 as the accuracy. Note that this result is better than the case of *maxDepth* = 0, which is a baseline, but there is significant difference in the accuracy between the results on validation and training data. This is probably because of the different distribution of values in the feature space.

References

- [1] Fayyad, Usama M. and Irani, Keki B. 1992. On the Handling of Continuous-Valued Attributes in Decision Tree Generation. *Machine Learning*, 8, pp. 87 - 102.